v2.0 · 2026

A FRAMEWORK FOR LLM RESEARCH AND AUDITS

PRISM

When one prompt isn't enough — and ad hoc multi-prompt work falls apart.

Prompts · Research · Iteration · Synthesis · Master

THE PAIN

By session three, nobody remembers the original thesis.

Serious research spans competitive, technical, financial, and strategic angles. One prompt goes shallow. Splitting helps — until the splits drift apart.

Outputs drift

Specialist passes contradict each other as sources stack up.

Sources collide

Two LLMs give different numbers for the same benchmark.

Context blows up

Session three is fighting for token headroom.

State is lost

“Wait — where did we leave off?” every time you resume.

THE THESIS

Serious research doesn't fit in one prompt.

A single LLM session either answers everything shallowly — or produces a long, unfocused mass that's hard to trust.

WHAT PRISM v2 IS

A method that keeps research coherent across prompts, sessions, and context limits.

Self-driving at Setup

You bring subject and goal. Setup runs as iterative refinement against the Lens Library — seven probes grade the draft strategy until coverage saturates.

Hardened per prompt

Every dispatch carries a triple contract: Envelope (inbound), Self-check (substrate verification), Output (outbound file with structured signature).

Multi-LLM by design

Vendor Selection chooses per dispatch. Triangulation reconciles divergence at convergence. Cross-vendor disagreement is signal, not noise.

Continuous state

The Master file carries every artifact across every session. “What’s next?” resumes any project — start, mid-flight, or after a long pause.

MOBILE-FIRST BY DESIGN

Designed around how you actually work.

The real substrate is a human moving artifacts between vendor chats, often on a phone. PRISM is built for that — file-based outputs, filename discipline, operator hints, and "What's next" all serve it directly.

Plain markdown, plain chat

One attached file, any vendor, any device. The framework runs in the same chat surface you already use. Manual paste-in is the default execution mode — not a fallback.

One file, infinite resume

The Master is a single .md you can email to yourself, save to a phone, share with a collaborator. State doesn't live in a service — it lives in the file.

Mobile-readable rendering

Release-notes use heading format, not tables. No unescaped pipes in code spans inside cells. Markdown that renders cleanly on a 6-inch screen — the framework dogfoods its own rules.

Three-leg design (§1.3): Operator · Substrate · Methodology — every mechanic earns its place by serving all three.

TWO SESSION TYPES

v2 names the split that v1.x ran inconsistently.

ORCHESTRATION

Claude with the framework attached

- Master state, Setup probes, Monitor fires, convergence reasoning.
- Substrate self-check at session open (SP-13).
- Closes by writing What's next to the Master.

EXECUTION

Any vendor, single dispatched prompt

- One prompt, one file out, no framework overhead.
- Loads only what the Envelope's Attachments field names.
- Vendor per the Envelope: Claude, Gemini, Perplexity, ChatGPT.

Crossing the boundary is the triple contract →

THE TRIPLE CONTRACT

The load-bearing interface between sessions.

Three blocks. Visually distinctive. Self-contained. Survives mode changes (paste-chat today, automated dispatch later) unchanged.

1

ENVELOPE · *Inbound*

Prompt ID, Master version, Vendor + config, Dispatch shape, Attachments, Expected output, Operator hints.

2

SELF-CHECK · *Substrate verification*

Vendor declares model identity, confirms match against Envelope's Vendor and Vendor config fields. Halts on mismatch.

3

OUTPUT · *Outbound*

A single .md file wrapped in PRISM Execution Output signature. Filename + next-action instruction baked in.

THE LENS LIBRARY

Coverage as a first-class artifact.

A reference catalog of audit-scope lenses graded against the draft Prompt Strategy. Probe 1 surfaces what a freehand scope would have silently missed.

5	18	7	3
Universal lenses	**Domain lenses**	**Setup probes**	**Layers of readiness**
Always evaluated	Across 6 packs	Grade vs. Library	Clear P0→P1

Each lens = material question · evidence class · specialist type. External frameworks (WCAG, OWASP) appear as provenance, not rubrics.

SETUP AS ITERATIVE REFINEMENT

P0.1 → P0.2 → … until coverage saturates.

Probe 1 grades the draft strategy against the Library each iteration. Probes 6 & 7 iterate early. Probes 2 & 4 iterate per turn. Probes 3 & 5 run once.

P1	P2	P3	P4	P5	P6	P7
Coverage grading	Adversarial Scope	Decision Framing	Pre-mortem	Falsifier	Domain Recon	User Voice

THREE-LAYER READINESS — P0 → P1 BOUNDARY

1 Structural completeness

Every planned prompt has all six required fields.

2 Library coverage saturation

Every applicable Lens covered or explicitly out of scope.

3 Operator ratification

Free-form approval. Silence is never consent (SP-9).

MASTER + WHAT'S NEXT

Continuous state. One question to resume.

Master: a single living markdown file holding every artifact — Decision brief, Stakeholder register, Claim inventory, Jurisdiction map, Prompt Strategy, Findings, Rerun Register, Learnings, Changelog.

ATTACH THE FILE. ASK ONE QUESTION.

"What's next?"

Universal entry point. Works for start, resume, or "where was I?" after a long pause.

You never reconstruct from memory. The Master is the source of truth. (SP-1)

CONTEXT-PRESSURE FRAMEWORK

The session knows when it’s slipping.

Seven signals — three volumetric, four behavioral — compounded qualitatively into four bands. Behavioral signals weight higher. Operator override is authoritative.

Comfortable

Normal operation

Getting warm

Continuous-curation posture

Curate now

Operator-elective migration

Migrate

Mandatory handoff

At Migrate — mandatory Migration handoff: a structured artifact, not pushing through a degraded session.

WHERE IT FITS

Any research problem too big for one prompt and too long for one session.

Product audits

Competitive · UX · technical · strategy

Competitive landscape

Market definition · segmentation · positioning

Market sizing

Multi-source cross-checked. Triangulation built in.

Investment due diligence

IC-ready memo with risk register

Technical DD

Engineering risk · regulatory · data lawfulness

Market entry

Go / no-go recommendation with falsifier

“Any research problem where a single prompt would either be too shallow or too long to trust.”

HOW TO START

Three steps.

1

Grab PRISM.md

From github.com/Ronkupper/PRISM. One file: framework + Lens Library + skill frontmatter.

2

Attach or install

Attach to a Claude session, or install as a Claude Skill so it auto-triggers on any project Master.

3

Send your subject

Hand Claude the project. Setup runs. Approve the strategy. Send “What’s next?”